

# ECE 637 Deep Learning Lab Exercises

Name: Alim Karimi

## Section 1

### Exercise 1.1

1. Create two lists, **A** and **B** : **A** contains 3 arbitrary numbers and **B** contains 3 arbitrary strings.
2. Concatenate two lists into a bigger list and name that list **C** .
3. Print the first element in **C** .
4. Print the second last element in **C** via negative indexing.
5. Remove the second element of **A** from **C** .
6. Print **C** again.

```
In [19]: # ----- YOUR CODE -----
import numpy as np

#1
A = [1,2,3]
B = ['hi', 'hello', 'bye']

#2 - concatenate lists
C = A + B

#3 print first element in C
print(C[0])

#4 - print second last element in C via negative indexing
print(C[-2])

#5- remove the second element of A from C:
rm = C.pop(-2)

#6 - print C again
print(C) #hello is not in the list now

1
hello
[1, 2, 3, 'hi', 'bye']
```

### Exercise 1.2

In this exercise, you will use a low-pass IIR filter to remove noise from a sine-wave signal.

You should organize your plots in a 3x1 subplot format.

1. Generate a discrete-time signal,  $x$ , by sampling a 2Hz continuous time sine wave signal with peak amplitude 1 from time 0s to 10s and at a sampling frequency of 500 Hz. Display the signal,  $x$ , from time 4s to 6s in the first row of a 3x1 subplot with the title "original signal".
2. Add Gaussian white random noise with 0 mean and standard deviation 0.1 to  $x$  and call it  $x_n$ . Display  $x_n$  from 4s to 6s on the second row of the subplot with the title "input signal".
3. Design a low-pass butterworth IIR filter of order 5 with a cut-off frequency of 4Hz, designed to filter out the noise. Hint: Use the `signal.butter` function and note that the frequencies are relative to the Nyquist frequency. Apply the IIR filter to  $x_n$ , and name the output  $y$ . Hint: Use `signal.filtfilt` function. Plot  $y$  from 4s to 6s on the third row of the subplot with the title "filtered signal".

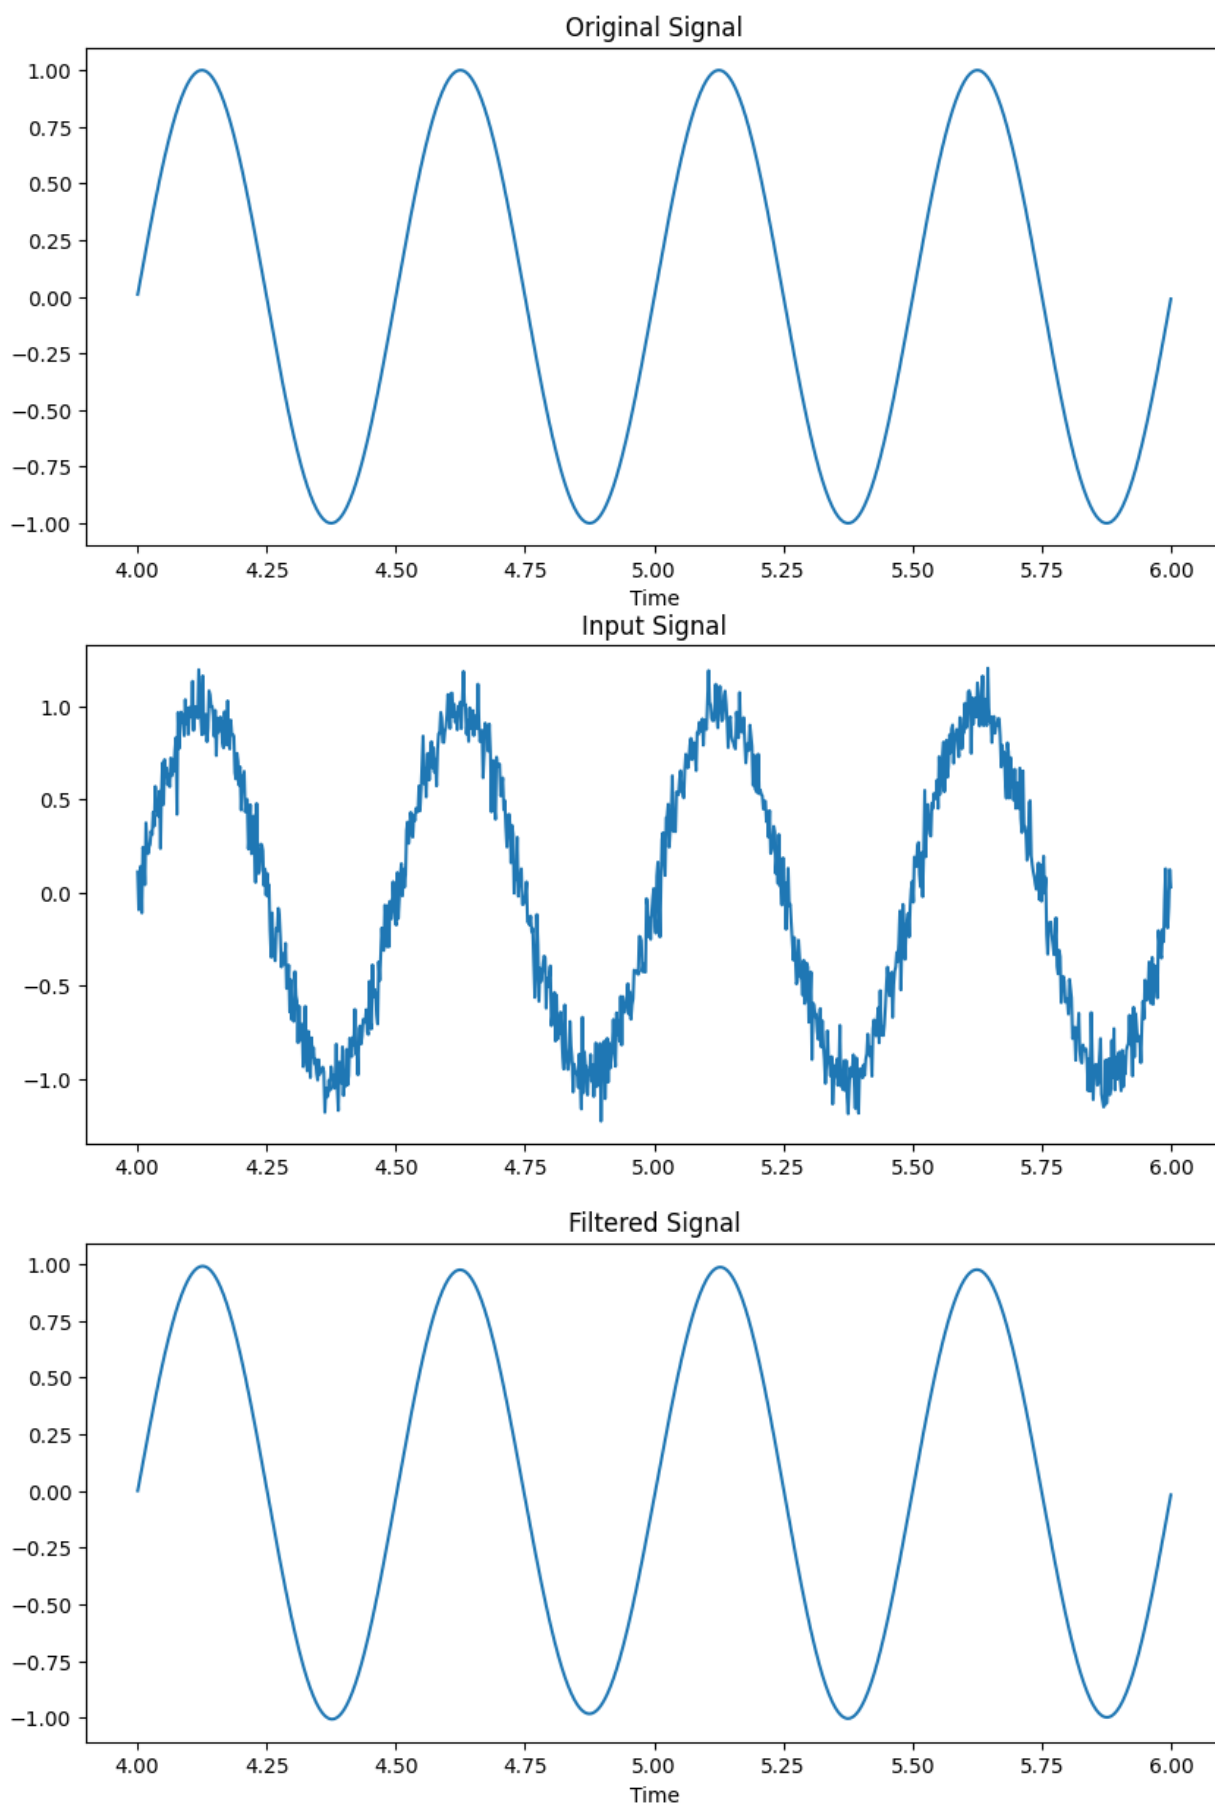
```
In [34]: import numpy as np                                # import the numpy packages and use a sh
import matplotlib.pyplot as plt                          # again import the matplotlib's pyplot p
from scipy import signal                                # import a minor package signal from sci
plt.figure(figsize=(10, 15))                            # fix the plot size

# ----- YOUR CODE -----
#part 1
t = np.linspace(0, 10, 500 * 10) #sampling freq of 500 Hz for 10 seconds means
x = np.sin(2 * np.pi * 2 * t)
plt.subplot(3,1,1)
plt.plot(t[2000:3000], x[2000:3000])
plt.xlabel('Time')

plt.title("Original Signal")

#part 2
mu = 0 #mean
sigma = 0.1 #std dev
noise = np.random.randn(500 * 10)*sigma + mu
x_n = x + noise
plt.subplot(3,1,2)
plt.plot(t[2000:3000], x_n[2000:3000])
plt.title("Input Signal")

#part 3
#Nyquist is half of the sampling rate (i.e, 0.5 * 500 = 250 Hz)
b, a = signal.butter(N = 5, Wn = 4, btype='low', analog=False, output='ba', fs=
#note: can also do b, a = signal.butter(N = 5, Wn = (4 / Nyquist_freq), btype=
#but in this case note that Wn is between 0 and 1 and relative to Nyquist
plt.subplot(3,1,3)
y = signal.filtfilt(b, a, x_n)
plt.plot(t[2000:3000], y[2000:3000])
plt.title("Filtered Signal")
plt.xlabel('Time')
plt.show()
```



## Section 2

### Exercise 2.1

- Plot the third image in the test data set
- Find the corresponding label for this image and make it the title of the figure

```
In [ ]: import keras
from keras.datasets import mnist
(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

train_images = train_images.reshape((60000, 28, 28, 1))
test_images = test_images.reshape((10000, 28, 28, 1))

# ----- YOUR CODE -----

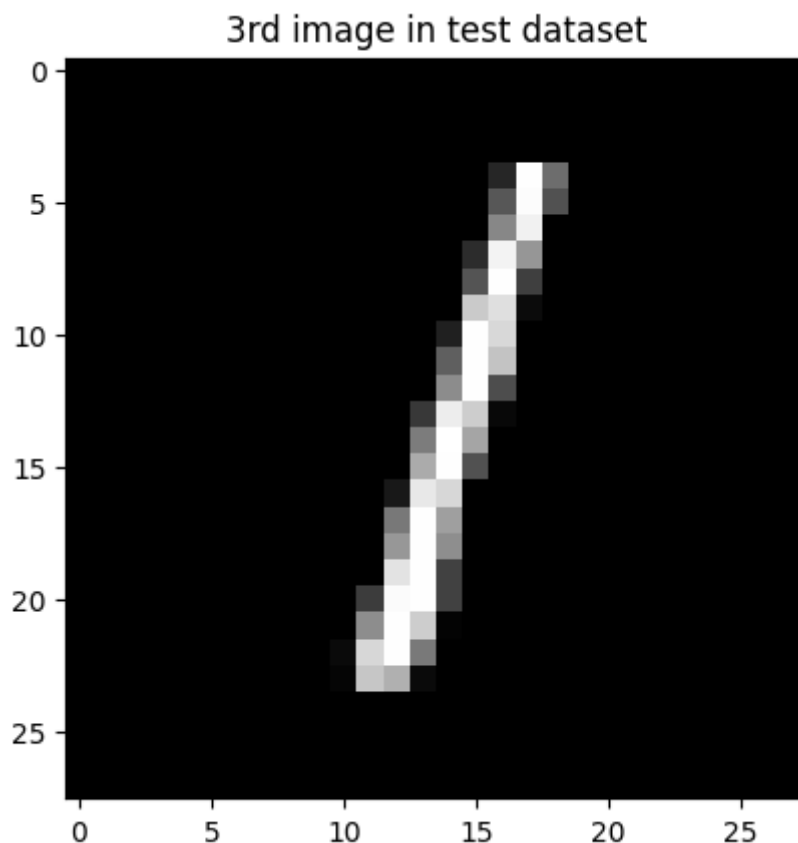
#Plot the 3rd image in the Test dataset
#print(test_images.shape)
plt.imshow(test_images[2], cmap='gray')
plt.title("3rd image in test dataset")

print("label for 3rd image in test dataset is:", test_labels[2])
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>

11490434/11490434 [=====] - 2s 0us/step

label for 3rd image in test dataset is: 1



## Exercise 2.2

It is usually helpful to have an accuracy plot as well as a loss value plot to get an intuitive sense of how effectively the model is being trained.

- Add code to this example for plotting two graphs with the following requirements:
  - Use a 1x2 subplot with the left subplot showing the loss function and right subplot showing the accuracy.
  - For each graph, plot the value with respect to epochs. Clearly label the x-axis, y-axis and the title.

(Hint: The value of loss and accuracy are stored in the `hist` variable. Try to print out `hist.history` and `hist.history.keys()`.)

```
In [35]: import keras
from keras.datasets import mnist
from keras import models
from keras import layers
from keras.utils import to_categorical

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

train_images = train_images.reshape((60000, 28, 28, 1))
test_images = test_images.reshape((10000, 28, 28, 1))

network = models.Sequential()
network.add(layers.Flatten(input_shape=(28, 28, 1))) #network takes in a 28*28
network.add(layers.Dense(512, activation='relu')) #takes the 784 --> 512, apply
network.add(layers.Dense(10, activation='softmax')) #takes the 512 --> 10, apply

network.summary()

network.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=[

train_images_nor = train_images.astype('float32') / 255 #rescale data to 0-1
test_images_nor = test_images.astype('float32') / 255 #rescale data to 0-1 (fro

train_labels_cat = to_categorical(train_labels) #one hot encoded label (i.e 3
test_labels_cat = to_categorical(test_labels)

hist = network.fit(train_images_nor, train_labels_cat, epochs=20, batch_size=12

test_loss, test_acc = network.evaluate(test_images_nor, test_labels_cat)
print('test_accuracy:', test_acc)
```

Model: "sequential\_6"

Layer (type)	Output Shape	Param #
flatten_4 (Flatten)	(None, 784)	0
dense_12 (Dense)	(None, 512)	401920
dense_13 (Dense)	(None, 10)	5130

=====  
 Total params: 407,050  
 Trainable params: 407,050  
 Non-trainable params: 0

Epoch 1/20  
 469/469 [=====] - 2s 4ms/step - loss: 0.2627 - accuracy: 0.9236  
 Epoch 2/20  
 469/469 [=====] - 2s 4ms/step - loss: 0.1077 - accuracy: 0.9680  
 Epoch 3/20  
 469/469 [=====] - 2s 3ms/step - loss: 0.0709 - accuracy: 0.9787  
 Epoch 4/20  
 469/469 [=====] - 2s 3ms/step - loss: 0.0515 - accuracy: 0.9843  
 Epoch 5/20  
 469/469 [=====] - 1s 3ms/step - loss: 0.0390 - accuracy: 0.9881  
 Epoch 6/20  
 469/469 [=====] - 2s 3ms/step - loss: 0.0290 - accuracy: 0.9915  
 Epoch 7/20  
 469/469 [=====] - 2s 3ms/step - loss: 0.0221 - accuracy: 0.9935  
 Epoch 8/20  
 469/469 [=====] - 2s 3ms/step - loss: 0.0166 - accuracy: 0.9955  
 Epoch 9/20  
 469/469 [=====] - 2s 4ms/step - loss: 0.0130 - accuracy: 0.9966  
 Epoch 10/20  
 469/469 [=====] - 2s 4ms/step - loss: 0.0094 - accuracy: 0.9979  
 Epoch 11/20  
 469/469 [=====] - 2s 3ms/step - loss: 0.0073 - accuracy: 0.9984  
 Epoch 12/20  
 469/469 [=====] - 2s 4ms/step - loss: 0.0050 - accuracy: 0.9990  
 Epoch 13/20  
 469/469 [=====] - 1s 3ms/step - loss: 0.0036 - accuracy: 0.9993  
 Epoch 14/20  
 469/469 [=====] - 2s 3ms/step - loss: 0.0024 - accuracy: 0.9997  
 Epoch 15/20  
 469/469 [=====] - 2s 5ms/step - loss: 0.0017 - accuracy: 0.9998

```

Epoch 16/20
469/469 [=====] - 4s 8ms/step - loss: 0.0011 - accuracy: 0.9999
Epoch 17/20
469/469 [=====] - 2s 5ms/step - loss: 6.5604e-04 - accuracy: 1.0000
Epoch 18/20
469/469 [=====] - 3s 6ms/step - loss: 4.9241e-04 - accuracy: 1.0000
Epoch 19/20
469/469 [=====] - 2s 4ms/step - loss: 3.8513e-04 - accuracy: 1.0000
Epoch 20/20
469/469 [=====] - 1s 3ms/step - loss: 3.2693e-04 - accuracy: 1.0000
313/313 [=====] - 1s 2ms/step - loss: 0.0672 - accuracy: 0.9829
test_accuracy: 0.9829000234603882

```

```

In [36]: import matplotlib.pyplot as plt

fig, ax = plt.subplots(2, figsize=(10,6))
fig.tight_layout(pad=5.0)

# ----- YOUR CODE -----
print(hist.history)
print(hist.history.keys())
ax[0].plot(hist.history['loss'])
ax[0].set_title("Loss")
ax[0].set_xlabel("Epochs")
ax[0].set_ylabel("Loss")

ax[1].plot(hist.history['accuracy'])
ax[1].set_title("Accuracy")
ax[1].set_xlabel("Epochs")
ax[1].set_ylabel("Accuracy")


plt.show()

```

```

{'loss': [0.26269108057022095, 0.10765278339385986, 0.07094603031873703, 0.051503609865903854, 0.038977112621068954, 0.02895783632993698, 0.02207002602517605, 0.016626844182610512, 0.012956601567566395, 0.009391698986291885, 0.0073214382864534855, 0.005036741495132446, 0.003593217581510544, 0.002376915654167533, 0.0017094099894165993, 0.0010693982476368546, 0.0006560396286658943, 0.000492409395519644, 0.00038513014442287385, 0.00032693208777345717], 'accuracy': [0.9235666394233704, 0.9679999947547913, 0.978683352470398, 0.9843000173568726, 0.988099992275238, 0.9914833307266235, 0.9934666752815247, 0.9954833388328552, 0.9966166615486145, 0.9978500008583069, 0.9983833432197571, 0.9989833235740662, 0.9992833137512207, 0.9996833205223083, 0.9997666478157043, 0.9999499917030334, 0.9999833106994629, 0.9999833106994629, 1.0, 1.0]}
dict_keys(['loss', 'accuracy'])

```



## Exercise 2.3

Use the dense network from Section 2 as the basis to construct of a deeper network with

- 5 dense hidden layers with dimensions [512, 256, 128, 64, 32] each of which uses a ReLU non-linearity

**Question:** Will the accuracy on the testing data always get better if we keep making the neural network larger?

No, it doesn't necessarily get better - it is very close to the accuracy of the smaller network, but actually a little bit less accurate in this case.

```
In [37]: import keras
from keras import models
from keras import layers

# ----- YOUR CODE -----
network = models.Sequential()
network.add(layers.Flatten(input_shape=(28, 28, 1))) #network takes in a 28*28
network.add(layers.Dense(512, activation='relu')) #takes the 784 --> 512, apply
network.add(layers.Dense(256, activation='relu'))
network.add(layers.Dense(128, activation='relu'))
network.add(layers.Dense(64, activation='relu'))
network.add(layers.Dense(32, activation='relu'))
network.add(layers.Dense(10, activation='softmax')) #takes the 512 --> 10, apply

network.summary()
```



Model: "sequential\_7"

Layer (type)	Output Shape	Param #
flatten_5 (Flatten)	(None, 784)	0
dense_14 (Dense)	(None, 512)	401920
dense_15 (Dense)	(None, 256)	131328
dense_16 (Dense)	(None, 128)	32896
dense_17 (Dense)	(None, 64)	8256
dense_18 (Dense)	(None, 32)	2080
dense_19 (Dense)	(None, 10)	330
Total params: 576,810		
Trainable params: 576,810		
Non-trainable params: 0		

```
In [38]: import keras
from keras.datasets import mnist
from keras.utils import to_categorical

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

train_images = train_images.reshape((60000, 28, 28, 1))
test_images = test_images.reshape((10000, 28, 28, 1))

network.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=[

train_images_nor = train_images.astype('float32') / 255
test_images_nor = test_images.astype('float32') / 255

train_labels_cat = to_categorical(train_labels)
test_labels_cat = to_categorical(test_labels)

hist = network.fit(train_images_nor, train_labels_cat, epochs=20, batch_size=12

test_loss, test_acc = network.evaluate(test_images_nor, test_labels_cat)
print('test accuracy:', test_acc)
```

```
Epoch 1/20
469/469 [=====] - 4s 4ms/step - loss: 0.2973 - accuracy: 0.9097
Epoch 2/20
469/469 [=====] - 2s 4ms/step - loss: 0.1035 - accuracy: 0.9686
Epoch 3/20
469/469 [=====] - 2s 4ms/step - loss: 0.0675 - accuracy: 0.9797
Epoch 4/20
469/469 [=====] - 3s 6ms/step - loss: 0.0517 - accuracy: 0.9841
Epoch 5/20
469/469 [=====] - 2s 4ms/step - loss: 0.0393 - accuracy: 0.9879
Epoch 6/20
469/469 [=====] - 2s 4ms/step - loss: 0.0289 - accuracy: 0.9913
Epoch 7/20
469/469 [=====] - 2s 4ms/step - loss: 0.0260 - accuracy: 0.9924
Epoch 8/20
469/469 [=====] - 2s 4ms/step - loss: 0.0217 - accuracy: 0.9936
Epoch 9/20
469/469 [=====] - 2s 4ms/step - loss: 0.0175 - accuracy: 0.9948
Epoch 10/20
469/469 [=====] - 2s 5ms/step - loss: 0.0158 - accuracy: 0.9949
Epoch 11/20
469/469 [=====] - 2s 5ms/step - loss: 0.0144 - accuracy: 0.9958
Epoch 12/20
469/469 [=====] - 2s 4ms/step - loss: 0.0112 - accuracy: 0.9966
Epoch 13/20
469/469 [=====] - 2s 4ms/step - loss: 0.0104 - accuracy: 0.9968
Epoch 14/20
469/469 [=====] - 2s 4ms/step - loss: 0.0084 - accuracy: 0.9975
Epoch 15/20
469/469 [=====] - 2s 4ms/step - loss: 0.0111 - accuracy: 0.9974
Epoch 16/20
469/469 [=====] - 2s 5ms/step - loss: 0.0075 - accuracy: 0.9981
Epoch 17/20
469/469 [=====] - 2s 5ms/step - loss: 0.0067 - accuracy: 0.9980
Epoch 18/20
469/469 [=====] - 2s 4ms/step - loss: 0.0073 - accuracy: 0.9978
Epoch 19/20
469/469 [=====] - 2s 4ms/step - loss: 0.0057 - accuracy: 0.9985
Epoch 20/20
469/469 [=====] - 2s 4ms/step - loss: 0.0063 - accuracy: 0.9984
```

```
313/313 [=====] - 1s 3ms/step - loss: 0.1180 - accuracy: 0.9841
test_accuracy: 0.9840999841690063
```

## Section 3

### Exercise 3.1

In this exercise, you will access the relationship between the feature extraction layer and classification layer. The example above uses two sets of convolutional layers and pooling layers in the feature extraction layer and two dense layers in the classification layers. The overall performance is around 98% for both training and test dataset. In this exercise, try to create a similar CNN network with the following requirements:

- Achieve the overall accuracy higher than 99% for training and testing dataset.
- Keep the total number of parameters used in the network lower than 100,000.

```
In [39]: import keras
from keras import models
from keras import layers

network = models.Sequential()

# ----- YOUR CODE -----

# ---- Feature extraction section
# First Layer
network.add(layers.Conv2D(16, (3, 3), activation='relu', input_shape=(28, 28, 1)))
network.add(layers.MaxPooling2D((2, 2)))
# Second Layer
network.add(layers.Conv2D(32, (3, 3), activation='relu'))
network.add(layers.MaxPooling2D((2, 2)))
#add a third layer:
# network.add(layers.Conv2D(64, (3,3), activation='relu'))
# network.add(layers.MaxPooling2D((2,2)))

# ---- Classification section
# Rearrange the data
network.add(layers.Flatten())
# Third Layer
#network.add(layers.Dense(512, activation='relu'))
#third layer
network.add(layers.Dense(64, activation = 'relu'))
# fifth Layer
network.add(layers.Dense(10, activation='softmax'))

network.summary()
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
=====		
conv2d_8 (Conv2D)	(None, 26, 26, 16)	160
max_pooling2d_4 (MaxPooling 2D)	(None, 13, 13, 16)	0
conv2d_9 (Conv2D)	(None, 11, 11, 32)	4640
max_pooling2d_5 (MaxPooling 2D)	(None, 5, 5, 32)	0
flatten_6 (Flatten)	(None, 800)	0
dense_20 (Dense)	(None, 64)	51264
dense_21 (Dense)	(None, 10)	650
=====		
Total params: 56,714		
Trainable params: 56,714		
Non-trainable params: 0		

```
In [ ]: from keras.datasets import mnist
        from keras.utils import to_categorical

        (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

        train_images = train_images.reshape((60000, 28, 28, 1))
        train_images_nor = train_images.astype('float32') / 255
        test_images = test_images.reshape((10000, 28, 28, 1))
        test_images_nor = test_images.astype('float32') / 255

        train_labels_cat = to_categorical(train_labels)
        test_labels_cat = to_categorical(test_labels)

        network.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=[
        network.fit(train_images_nor, train_labels_cat, epochs=6, batch_size=64)

        test_loss, test_acc = network.evaluate(test_images_nor, test_labels_cat)
        print('test_accuracy:', test_acc)
```

```

Epoch 1/6
938/938 [=====] - 10s 4ms/step - loss: 0.1955 - accuracy: 0.9421
Epoch 2/6
938/938 [=====] - 4s 4ms/step - loss: 0.0585 - accuracy: 0.9817
Epoch 3/6
938/938 [=====] - 4s 4ms/step - loss: 0.0418 - accuracy: 0.9875
Epoch 4/6
938/938 [=====] - 4s 4ms/step - loss: 0.0323 - accuracy: 0.9904
Epoch 5/6
938/938 [=====] - 4s 4ms/step - loss: 0.0266 - accuracy: 0.9918
Epoch 6/6
938/938 [=====] - 4s 5ms/step - loss: 0.0228 - accuracy: 0.9933
313/313 [=====] - 1s 3ms/step - loss: 0.0363 - accuracy: 0.9875
test_accuracy: 0.987500011920929

```

## Section 4

### Exercise 4.1

In this exercise you will need to create the entire neural network that does image denoising tasks. Try to mimic the code provided above and follow the structure as provided in the instructions below.

#### Task 1: Create the datasets

1. Import necessary packages
2. Load the MNIST data from Keras, and save the training dataset images as `train_images`, save the test dataset images as `test_images`
3. Add additive white gaussian noise to the train images as well as the test images and save the noisy images to `train_images_noisy` and `test_images_noisy` respectively. The noise should have mean value 0, and standard deviation 0.4. (Hint: Use [np.random.normal](#))
4. Show the first image in the training dataset as well as the test dataset (plot the images in 1 x 2 subplot form)

```

In [ ]: # ----- YOUR CODE -----
        ### Part 1, load packages:
        from keras.datasets import mnist
        from keras.utils import to_categorical

        #part 2, load MNIST data from Keras, save the training dataset as train_images,
        (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

```

```

train_images = train_images.reshape((60000, 28, 28, 1))

train_images_nor = train_images.astype('float32') / 255
test_images = test_images.reshape((10000, 28, 28, 1))
test_images_nor = test_images.astype('float32') / 255

#Part 3, add additive white gaussian noise to the train images and test images.
noise_train = np.random.normal(loc=0, scale=0.4, size = (60000,28,28,1))
noise_test = np.random.normal(loc = 0, scale = 0.4, size = (10000, 28, 28, 1))
train_images_noisy = train_images_nor + noise_train
test_images_noisy = test_images_nor + noise_test


#part 4, show the first image in training and first image in test dataset

plt.subplot(1,2,1)
plt.imshow(train_images_noisy[0], cmap='gray')
plt.title("noisy train image")

plt.subplot(1,2,2)
plt.imshow(test_images_noisy[0], cmap='gray')
plt.title("noisy test image")

```

Out[ ]: Text(0.5, 1.0, 'noisy test image')



## Task 2: Create the neural network model

1. Create a sequential model called `encoder` with the following layers sequentially:
  - convolutional layer with 32 output channels, 3x3 kernel size, and the padding convention 'same' with 'relu' activation function.
  - max pooling layer with 2x2 kernel size
  - convolutional layer with 16 output channels, 3x3 kernel size, and the padding convention 'same' with 'relu' activation function.
  - max pooling layer with 2x2 kernel size
  - convolutional layer with 8 output channels, 3x3 kernel size, and the padding convention 'same' with 'relu' activation function and name the layer as 'convOutput'.
  - flatten layer

- dense layer with output dimension as `encoding_dim` with `'relu'` activation function.
2. Create a sequential model called `decoder` with the following layers sequentially:
- dense layer with the input dimension as `encoding_dim` and the output dimension as the product of the output dimensions of the `'convOutput'` layer.
  - reshape layer that convert the tensor into the same shape as `'convOutput'`
  - convolutional layer with `8` output channels, `3x3` kernel size, and the padding convention `'same'` with `'relu'` activation function.
  - upsampling layer with `2x2` kernel size
  - convolutional layer with `16` output channels, `3x3` kernel size, and the padding convention `'same'` with `'relu'` activation function.
  - upsampling layer with `2x2` kernel size
  - convolutional layer with `32` output channels, `3x3` kernel size, and the padding convention `'same'` with `'relu'` activation function
  - convolutional layer with `1` output channels, `3x3` kernel size, and the padding convention `'same'` with `'sigmoid'` activation function
3. Create a sequential model called `autoencoder` with the following layers sequentially:
- `encoder` model
  - `decoder` model

```
In [ ]: # ----- YOUR CODE -----
# Build Encoder
encoder = models.Sequential()
encoder.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=(7, 7, 3)))
encoder.add(layers.MaxPooling2D((2, 2), padding='same'))
encoder.add(layers.Conv2D(16, (3, 3), activation='relu', padding='same'))
encoder.add(layers.MaxPooling2D((2, 2), padding='same'))
encoder.add(layers.Conv2D(8, (3, 3), activation='relu', padding='same', name='convOutput'))
encoder.add(layers.Flatten())
encoding_dim = 32
encoder.add(layers.Dense(encoding_dim, activation='relu'))

# shape considerations
convShape = encoder.get_layer('convOutput').output_shape[1:]
print(convShape)
denseShape = convShape[0]*convShape[1]*convShape[2]
print(denseShape)
#Build Decoder
decoder = models.Sequential()
decoder.add(layers.Dense(denseShape, input_shape=(encoding_dim,)))
decoder.add(layers.Reshape(convShape))

decoder.add(layers.Conv2D(8, (3, 3), activation='relu', padding='same'))
decoder.add(layers.UpSampling2D((2, 2)))
decoder.add(layers.Conv2D(16, (3, 3), activation='relu', padding='same'))
decoder.add(layers.UpSampling2D((2, 2)))
decoder.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same'))
decoder.add(layers.Conv2D(1, (3, 3), activation='sigmoid', padding='same'))

(7, 7, 8)
392
```

```
In [ ]: encoder.summary()
        decoder.summary()

autoencoder = models.Sequential()
autoencoder.add(encoder)
autoencoder.add(decoder)
autoencoder.summary()
```



Model: "sequential\_3"

Layer (type)	Output Shape	Param #
conv2d_2 (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d_2 (MaxPooling 2D)	(None, 14, 14, 32)	0
conv2d_3 (Conv2D)	(None, 14, 14, 16)	4624
max_pooling2d_3 (MaxPooling 2D)	(None, 7, 7, 16)	0
convOutput (Conv2D)	(None, 7, 7, 8)	1160
flatten_3 (Flatten)	(None, 392)	0
dense_10 (Dense)	(None, 32)	12576

=====  
Total params: 18,680  
Trainable params: 18,680  
Non-trainable params: 0

Model: "sequential\_4"

Layer (type)	Output Shape	Param #
dense_11 (Dense)	(None, 392)	12936
reshape (Reshape)	(None, 7, 7, 8)	0
conv2d_4 (Conv2D)	(None, 7, 7, 8)	584
up_sampling2d (UpSampling2D )	(None, 14, 14, 8)	0
conv2d_5 (Conv2D)	(None, 14, 14, 16)	1168
up_sampling2d_1 (UpSampling 2D)	(None, 28, 28, 16)	0
conv2d_6 (Conv2D)	(None, 28, 28, 32)	4640
conv2d_7 (Conv2D)	(None, 28, 28, 1)	289

=====  
Total params: 19,617  
Trainable params: 19,617  
Non-trainable params: 0

Model: "sequential\_5"

Layer (type)	Output Shape	Param #
sequential_3 (Sequential)	(None, 32)	18680
sequential_4 (Sequential)	(None, 28, 28, 1)	19617

```
=====
Total params: 38,297
Trainable params: 38,297
Non-trainable params: 0
=====
```

---

### Task 3: Create the neural network model

Fit the model to the training data using the following hyper-parameters:

- `adam` optimizer
- `binary_crossentropy` loss function
- `20` training epochs
- batch size as `256`
- set `shuffle` as `True`

Compile the model and fit ...

```
In [ ]: autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
        history = autoencoder.fit(train_images_noisy, train_images_nor,
                                epochs=20,
                                batch_size=256,
                                shuffle=True)
```

```

Epoch 1/20
235/235 [=====] - 8s 12ms/step - loss: 0.1682
Epoch 2/20
235/235 [=====] - 2s 10ms/step - loss: 0.1385
Epoch 3/20
235/235 [=====] - 2s 10ms/step - loss: 0.1291
Epoch 4/20
235/235 [=====] - 2s 10ms/step - loss: 0.1236
Epoch 5/20
235/235 [=====] - 2s 10ms/step - loss: 0.1202
Epoch 6/20
235/235 [=====] - 3s 11ms/step - loss: 0.1177
Epoch 7/20
235/235 [=====] - 3s 12ms/step - loss: 0.1159
Epoch 8/20
235/235 [=====] - 2s 10ms/step - loss: 0.1143
Epoch 9/20
235/235 [=====] - 2s 10ms/step - loss: 0.1132
Epoch 10/20
235/235 [=====] - 3s 11ms/step - loss: 0.1120
Epoch 11/20
235/235 [=====] - 3s 14ms/step - loss: 0.1111
Epoch 12/20
235/235 [=====] - 2s 10ms/step - loss: 0.1102
Epoch 13/20
235/235 [=====] - 2s 10ms/step - loss: 0.1095
Epoch 14/20
235/235 [=====] - 2s 10ms/step - loss: 0.1088
Epoch 15/20
235/235 [=====] - 3s 11ms/step - loss: 0.1083
Epoch 16/20
235/235 [=====] - 3s 11ms/step - loss: 0.1076
Epoch 17/20
235/235 [=====] - 2s 10ms/step - loss: 0.1071
Epoch 18/20
235/235 [=====] - 2s 10ms/step - loss: 0.1067
Epoch 19/20
235/235 [=====] - 2s 10ms/step - loss: 0.1063
Epoch 20/20
235/235 [=====] - 3s 11ms/step - loss: 0.1057

```

**Task 4:** Create the neural network model (No need to write code, just run the following commands)

```

In [ ]: def showImages(input_imgs, encoded_imgs, output_imgs, size=1.5, groundTruth=None):

    numCols = 3 if groundTruth is None else 4

    num_images = input_imgs.shape[0]

    encoded_imgs = encoded_imgs.reshape((num_images, 1, -1))

    plt.figure(figsize=((numCols+encoded_imgs.shape[2])/input_imgs.shape[2])*size)

    pltIdx = 0
    col = 0
    for i in range(0, num_images):

```

```

        col += 1
    # plot input image
    pltIdx += 1
    ax = plt.subplot(num_images, numCols, pltIdx)
    plt.imshow(input_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    if col == 1:
        plt.title('Input Image')

    # plot encoding
    pltIdx += 1
    ax = plt.subplot(num_images, numCols, pltIdx)
    plt.imshow(encoded_imgs[i])
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    if col == 1:
        plt.title('Encoded Image')

    # plot reconstructed image
    pltIdx += 1
    ax = plt.subplot(num_images, numCols, pltIdx)
    plt.imshow(output_imgs[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
    if col == 1:
        plt.title('Reconstructed Image')

    if numCols == 4:
        # plot ground truth image
        pltIdx += 1
        ax = plt.subplot(num_images, numCols, pltIdx)
        plt.imshow(groundTruth[i].reshape(28, 28))
        plt.gray()
        ax.get_xaxis().set_visible(False)
        ax.get_yaxis().set_visible(False)

        if col == 1:
            plt.title('Ground Truth')

plt.show()

```

```

In [ ]: num_images = 10

input_labels = test_labels[0:num_images]
I = np.argsort(input_labels)

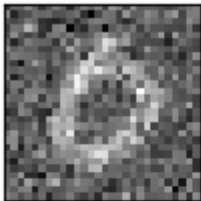
input_imgs = test_images_noisy[I]

encoded_imgs = encoder.predict(test_images_noisy[I])
output_imgs = decoder.predict(encoded_imgs)

showImages(input_imgs, encoded_imgs, output_imgs, size=1.5, groundTruth=test_in

1/1 [=====] - 0s 18ms/step
1/1 [=====] - 0s 18ms/step

```

Input Image		Reconstructed Image	Ground Truth
	Encoded Image 